

Assignment 9

Maksim Nikiforov

10/19/2021

Contents

Part 1: kNN	1
1.	1
2.	1
3.	2
4.	2
5.	2
Part 2: Ensemble	3
1.	3
2.	4
3.	4
4.	5

Part 1: kNN

1.

Read in the `heart.csv` data file.

We remove the `ST_Slope` column as instructed.

```
# Read in space-separated data using "read_csv()" function.
# Remove column ST_Slope, convert HeartDisease to a factor
heartData <- read.csv(file = "./heart.csv") %>% select(-ST_Slope)
```

2.

Create dummy columns corresponding to the values of `Sex`, `ChestPainType`, and `RestingECG` for use in our kNN fit.

We can follow the example from section 3.1 in the provided caret vignette to create dummy variables. `dummyVars` ignores integers, so we specify the full set of variables with the formula `~ .` and create a new data frame, `newHeartData`.

```
# Use dummyVars() and predict() to create new columns.
dummies <- dummyVars( ~ ., data = heartData)
newHeartData <- predict(dummies, newdata = heartData)
newHeartData <- as_data_frame(newHeartData)
```

```
## Warning: `as_data_frame()` was deprecated in tibble 2.0.0.
## Please use `as_tibble()` instead.
## The signature and semantics have changed, see `?as_tibble`.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.
```

```
newHeartData$HeartDisease <- as.factor(newHeartData$HeartDisease)
```

3.

Now split the data set you've created into a training and testing set. Use $p = 0.8$.

We can create train and test partitions with the `createDataPartition` function.

```
# p = 0.8 places 80% of observations in the training set
set.seed(50)
heartIndex <- createDataPartition(newHeartData$HeartDisease, p = 0.8, list = FALSE)
heartTrain <- newHeartData[heartIndex, ]
heartTest <- newHeartData[-heartIndex, ]
```

4.

Finally, train the kNN model. Use repeated 10 fold cross-validation, with the number of repeats being 3. You should also preprocess the data by centering and scaling. Lastly, set the `tuneGrid` so that you are considering values of k of 1, 2, 3, ..., 40.

The `caret` package allows us to pre-process data, fit and tune models on the training set, and predict on the test set.

```
#
kNNFit <- train(HeartDisease ~ ., data = heartTrain,
               method = "knn",
               preProcess = c("center", "scale"),
               trControl = trainControl(method = "repeatedcv",
                                       number = 10, repeats = 3),
               tuneGrid = data.frame(k = 1:40))
```

5.

Check how well your model does on the test set using the `confusionMatrix()` function.

```
# See confusion matrix on test set (correct/incorrect predictions)
confusionMatrix(data = heartTest$HeartDisease, reference = predict(kNNFit, newdata = heartTest))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 60 22
##           1 18 83
##
##           Accuracy : 0.7814
##           95% CI : (0.7145, 0.839)
##           No Information Rate : 0.5738
##           P-Value [Acc > NIR] : 2.949e-09
##
##           Kappa : 0.556
##
##           Mcnemar's Test P-Value : 0.6353
##
##           Sensitivity : 0.7692
##           Specificity : 0.7905
##           Pos Pred Value : 0.7317
```

```
##          Neg Pred Value : 0.8218
##          Prevalence : 0.4262
##          Detection Rate : 0.3279
##          Detection Prevalence : 0.4481
##          Balanced Accuracy : 0.7799
##
##          'Positive' Class : 0
##
```

Part 2: Ensemble

We'll look at predicting the same heart disease variable in this section as well, just instead of using kNN we'll use the following methods:

1.

A classification tree (use `method = rpart`: tuning parameter is `cp`, use values 0, 0.001, 0.002, ..., 0.1).

We use `seq()` to generate values from 0 to 0.1. With this method, we see an accuracy of approximately 80.33%.

```
# Create tuning parameter by generating regular sequence
rpartParam <- seq(from = 0, to = 0.1, by = 0.001)

# Classification tree fit
rpartFit <- train(HeartDisease ~ ., data = heartTrain,
  method = "rpart",
  preProcess = c("center", "scale"),
  trControl = trainControl(method = "repeatedcv",
    number = 10, repeats = 3),
  tuneGrid = data.frame(cp = rpartParam))

# See confusion matrix on test set (correct/incorrect predictions)
confusionMatrix(data = heartTest$HeartDisease, reference = predict(rpartFit, newdata = heartTest))

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0  1
##          0 57 25
##          1 11 90
##
##          Accuracy : 0.8033
##          95% CI : (0.7382, 0.8583)
##          No Information Rate : 0.6284
##          P-Value [Acc > NIR] : 2.307e-07
##
##          Kappa : 0.5958
##
##          Mcnemar's Test P-Value : 0.03026
##
##          Sensitivity : 0.8382
##          Specificity : 0.7826
##          Pos Pred Value : 0.6951
##          Neg Pred Value : 0.8911
##          Prevalence : 0.3716
```

```
##          Detection Rate : 0.3115
##    Detection Prevalence : 0.4481
##      Balanced Accuracy : 0.8104
##
##      'Positive' Class : 0
##
```

2.

A bagged tree (use `method = treebag`: no tuning parameter).

Here, we see a slightly lower accuracy of 79.23%.

```
# Bagged tree fit
baggedTreeFit <- train(HeartDisease ~ ., data = heartTrain,
  method = "treebag",
  preProcess = c("center", "scale"),
  trControl = trainControl(method = "repeatedcv",
    number = 10, repeats = 3))

# See confusion matrix on test set (correct/incorrect predictions)
confusionMatrix(data = heartTest$HeartDisease,
  reference = predict(baggedTreeFit, newdata = heartTest))
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0  1
##          0 57 25
##          1 13 88
##
##          Accuracy : 0.7923
##          95% CI : (0.7263, 0.8487)
##    No Information Rate : 0.6175
##    P-Value [Acc > NIR] : 2.954e-07
##
##          Kappa : 0.5743
##
## Mcnemar's Test P-Value : 0.07435
##
##          Sensitivity : 0.8143
##          Specificity : 0.7788
##          Pos Pred Value : 0.6951
##          Neg Pred Value : 0.8713
##          Prevalence : 0.3825
##          Detection Rate : 0.3115
##    Detection Prevalence : 0.4481
##          Balanced Accuracy : 0.7965
##
##      'Positive' Class : 0
##
```

3.

A random forest (use `method = rf`: tuning parameter is `mtry`, use vales of 1, 2, ..., 15).

Here, we see an accuracy of ~80.87%, on part with the classification tree method.

```
# Random forest fit
rfFit <- train(HeartDisease ~ ., data = heartTrain,
              method = "rf",
              preProcess = c("center", "scale"),
              trControl = trainControl(method = "repeatedcv",
                                       number = 10, repeats = 3),
              tuneGrid = data.frame(mtry = 1:15))

# See confusion matrix on test set (correct/incorrect predictions)
confusionMatrix(data = heartTest$HeartDisease,
                reference = predict(rfFit, newdata = heartTest))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 59 23
##           1 13 88
##
##           Accuracy : 0.8033
##           95% CI : (0.7382, 0.8583)
##           No Information Rate : 0.6066
##           P-Value [Acc > NIR] : 9.382e-09
##
##           Kappa : 0.5977
##
##           Mcnemar's Test P-Value : 0.1336
##
##           Sensitivity : 0.8194
##           Specificity : 0.7928
##           Pos Pred Value : 0.7195
##           Neg Pred Value : 0.8713
##           Prevalence : 0.3934
##           Detection Rate : 0.3224
##           Detection Prevalence : 0.4481
##           Balanced Accuracy : 0.8061
##
##           'Positive' Class : 0
##
```

4.

A boosted tree (use `method = gbm`: tuning parameters are `n.trees`, `interaction.depth`, `shrinkage`, and `n.minobsinnode`, use all combinations of `n.trees` of 25, 50, 100, 150, and 200, `interaction.depth` of 1, 2, 3, 4, `shrinkage = 0.1`, and `n.minobsinnode = 10`; Hint: use `expand.grid()` to create your data frame for `tuneGrid`).

With this method, our accuracy with the test data set is ~79.78%. B

```
# See confusion matrix on test set (correct/incorrect predictions)
confusionMatrix(data = heartTest$HeartDisease,
                reference = predict(boostedTreeFit, newdata = heartTest))
```

```
## Confusion Matrix and Statistics
```

```

##
##           Reference
## Prediction  0  1
##           0 58 24
##           1 15 86
##
##           Accuracy : 0.7869
##           95% CI : (0.7204, 0.8438)
##           No Information Rate : 0.6011
##           P-Value [Acc > NIR] : 7.088e-08
##
##           Kappa : 0.5646
##
## Mcnemar's Test P-Value : 0.2002
##
##           Sensitivity : 0.7945
##           Specificity : 0.7818
##           Pos Pred Value : 0.7073
##           Neg Pred Value : 0.8515
##           Prevalence : 0.3989
##           Detection Rate : 0.3169
##           Detection Prevalence : 0.4481
##           Balanced Accuracy : 0.7882
##
##           'Positive' Class : 0
##

```